



Chapter 2

Application Layer

Courtesy to the textbooks' authors and Pearson Addison-Wesley because many slides are adapted from the following textbooks and their associated slides.



Jim Kurose, Keith Ross, "Computer Networking: A Top Down Approach", 7th Edition, Pearson, 2016.



Jim Kurose, Keith Ross, "Computer Networking: A Top Down Approach", 8th Edition, Pearson, 2020.

All material copyright 1996-2020 J.F Kurose and K.W. Ross, All Rights Reserved

Application layer: overview

- Principles of network applications
- Web and HTTP
- E-mail, SMTP, IMAP
- The Domain Name System DNS
- P2P applications
- video streaming and content distribution networks
- socket programming with UDP and TCP
 - reading assignment

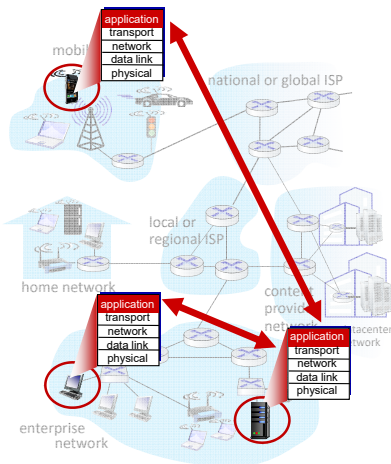
Where does a network app reside?

write programs that:

- run on (different) end systems
- communicate over network
- e.g., web server software communicates with browser software

no need to write software for network-core devices

- network-core devices do not run user applications
- applications on end systems allows for rapid app development, propagation




Client-server paradigm

server:

- always-on host
- permanent IP address
- often in data centers, for scaling

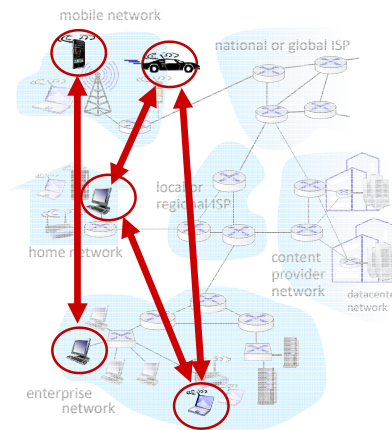
clients:

- contact, communicate with server
- may be intermittently connected
- may have dynamic IP addresses
- do *not* communicate directly with each other
- examples: HTTP, IMAP, FTP



Peer-peer architecture

- *no* always-on server
- arbitrary end systems directly communicate
- a peer *i*) requests service from other peers and *ii*) provides service in return to other peers
 - *self scalability* – new peers bring new service capacity, as well as new service demands
- peers are intermittently connected and change IP addresses
 - complex management
- example: P2P file sharing



5

How does two processes in hosts communicate?

process: program running within a host

- within same host, two processes communicate using **inter-process communication** (defined by OS)
- in different hosts, network app processes communicate by exchanging **messages** via **socket**

clients, servers

client process: process that initiates communication


server process: process that waits to be contacted

- note: an application with P2P architectures have both client process & server process

6

Sockets

- process sends/receives messages to/from its **socket**
- socket analogous to door (or mailbox)
 - sending process shoves message out door
 - sending process relies on transport infrastructure on other side of door to deliver message to socket at receiving process
 - two sockets involved: one on each side



7

Addressing processes

- to receive messages, process must have **identifier**
- host device has unique 32-bit IP address
- **Q:** does IP address of host on which process runs suffice for identifying the process?
 - **A:** no, *many* processes can be running on same host
- **identifier** includes both **IP address** and **port number** associated with process on host.
- example port numbers:
 - HTTP server: 80
 - mail server: 25
- to send HTTP message to www.nthu.edu.tw web server:
 - **IP address:** 140.114.69.135
 - **port number:** 80
- more shortly...

8

An application-layer protocol defines:

- **types of messages exchanged**
 - e.g., request, response, ...
- **message syntax**
 - what fields in messages & how fields are delineated
- **message semantics**
 - meaning of information in fields
- **rules** for when and how processes send & respond to messages

9

What transport service does an app need?

- data integrity
 - some apps (e.g., file transfer, web transactions) require 100% reliable data transfer
 - other apps (e.g., audio) can tolerate some loss
- timing
 - some apps (e.g., Internet telephony, interactive games) require low delay
- throughput
 - some apps (e.g., multimedia) require minimum amount of throughput
 - other apps (“elastic apps”) make use of whatever throughput they get
- security
 - encryption, data integrity, ...

10

Transport service requirements: common apps

| application | data loss | throughput | time sensitive? |
|------------------------|---------------|---|-----------------|
| file transfer/download | no loss | elastic | no |
| e-mail | no loss | elastic | no |
| Web documents | no loss | elastic | no |
| real-time audio/video | loss-tolerant | audio: 5Kbps-1Mbps video: 10Kbps-5Mbps | yes, 100's msec |
| streaming audio/video | loss-tolerant | same as above | yes, few secs |
| interactive games | loss-tolerant | Kbps+ | yes, 100's msec |
| text messaging | no loss | elastic | yes and no |

11

Internet transport protocols services

TCP service:

- **reliable transport** between sending and receiving process
- **flow control**: sender won't overwhelm receiver
- **congestion control**: throttle sender when network is overloaded
- **connection-oriented**: setup required between client and server processes
- **does not provide**: timing, minimum throughput guarantee, security

UDP service:

- **unreliable data transfer** between sending and receiving process
- **does not provide**: reliability, flow control, congestion control, timing, throughput guarantee, security, or connection setup.

Why is there UDP?

- is no-frills and lightweight
- provides minimal services

12

Internet applications, and transport protocols

| application | application layer protocol | transport protocol |
|------------------------|--|--------------------|
| file transfer/download | FTP [RFC 959] | TCP |
| e-mail | SMTP [RFC 5321] | TCP |
| Web documents | HTTP 1.1 [RFC 2616] | TCP |
| Internet telephony | SIP [RFC 3261], RTP [RFC 3550], or proprietary | TCP or UDP |
| streaming audio/video | HTTP [RFC 7320], DASH | TCP |

13

Application layer: overview

- Principles of network applications
- Web and HTTP
- E-mail, SMTP, IMAP
- The Domain Name System DNS
- P2P applications
- video streaming and content distribution networks
- socket programming with UDP and TCP



14

Web and HTTP

web page (or webpage)

- web page consists of *objects*
 - object is a file such as a HTML file, image, java applet, video clip, ...
 - objects can be stored on different web servers
- most web pages consist of a *base HTML-file* and several *referenced objects*

- each* addressable by a *URL* (Uniform Resource Locator), e.g.,
`www.someschool.edu/somefolder/index.html`

`www.someschool.edu/somefolder/index.html`
└──────────┘ └──────────┘
host name path name

`www.someschool.edu/somefolder/pic.jpg`

15

HTTP overview

HTTP: hypertext transfer protocol

- Web's application-layer protocol
- client/server model:
 - client*: browser that requests, receives, (using HTTP protocol), and displays Web objects
 - server*: Web server sends (using HTTP protocol) objects in response to requests



16

HTTP overview (continued)

HTTP uses TCP:

- client initiates TCP connection (creates socket) to server, port 80
- server accepts TCP connection from client
- HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
- TCP connection closed

HTTP is “stateless”

- server maintains *no* information about past client requests
- other mechanisms take care of “state”

17

HTTP connections: two types

Non-persistent HTTP

1. TCP connection opened
2. at most one object sent over TCP connection
3. TCP connection closed

downloading multiple objects required multiple connections

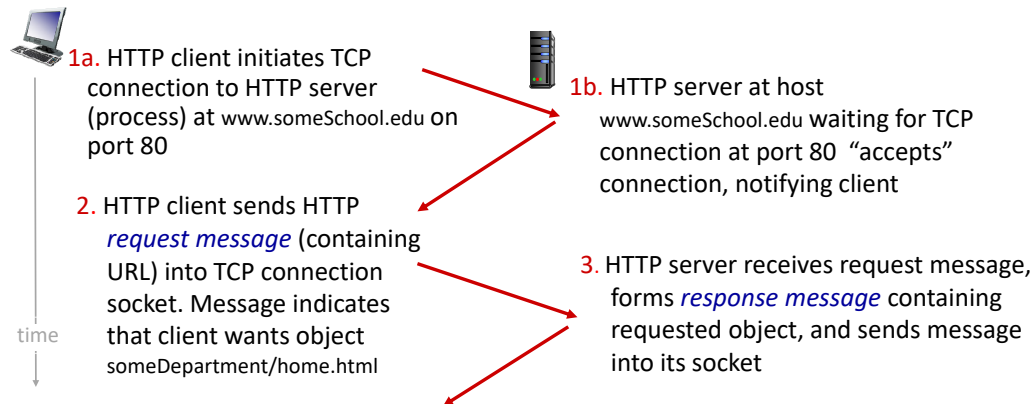
Persistent HTTP

- TCP connection opened to a server
- multiple objects can be sent over *single* TCP connection between client, and that server
- TCP connection closed

18

Non-persistent HTTP: example

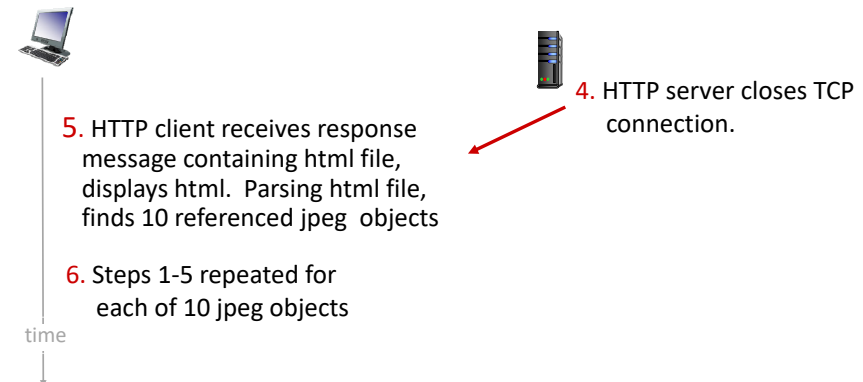
User enters URL: `www.someSchool.edu/someDepartment/home.html`
(containing text, references to 10 jpeg images)



19

Non-persistent HTTP: example (cont.)

User enters URL: `www.someSchool.edu/someDepartment/homepage.html`
(containing text, references to 10 jpeg images)




20

Non-persistent HTTP: response time

RTT (definition): time for a small packet to travel from client to server and back

HTTP response time (per object):

- one RTT to initiate TCP connection
- one RTT for HTTP request and first few bytes of HTTP response to return
- object/file transmission time



Non-persistent HTTP response time = 2RTT + file transmission time

21

Persistent HTTP (HTTP 1.1)

Non-persistent HTTP issues:

- requires 2 RTTs per object
- browsers often open multiple parallel TCP connections to fetch referenced objects in parallel
 - OS overhead for *each* TCP connection

Persistent HTTP (HTTP1.1):

- server leaves connection open after sending response
- subsequent HTTP messages between same client/server sent over open connection
- client sends requests as soon as it encounters a referenced object
- as little as one RTT for all the referenced objects (cutting response time in half)

22

HTTP/2

Key goal: decreased delay in multi-object HTTP requests

HTTP1.1: introduced **multiple, pipelined GETs** over single TCP connection

- server responds *in-order* (FCFS: first-come-first-served scheduling) to GET requests
- with FCFS, small object may have to wait for transmission (**head-of-line [HOL] blocking**) behind large object(s)
- loss recovery (retransmitting lost TCP segments) stalls object transmission

multiple parallel TCP connections: one for a single object

- drawback: high overhead (# of sockets maintained at servers), unfair

23

HTTP/2

Key goal: decreased delay in multi-object HTTP requests

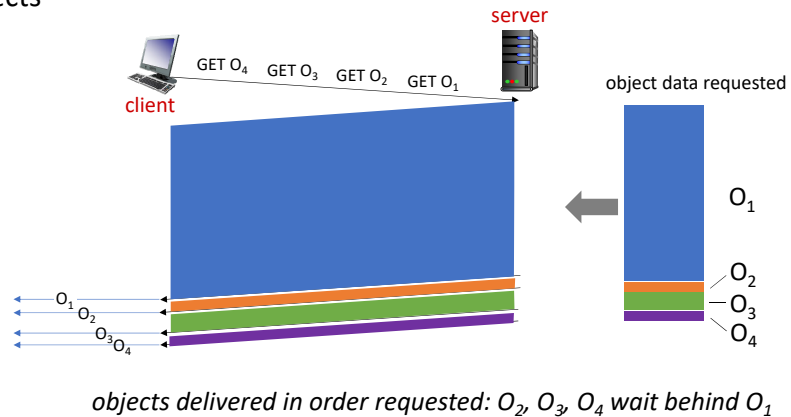
HTTP/2: [RFC 7540, 2015] increased flexibility at *server* in sending objects to client:

- methods, status codes, most header fields unchanged from HTTP 1.1
- transmission order of requested objects based on client-specified object priority (not necessarily FCFS)
- *push* unrequested objects to client
- divide objects into frames, schedule frames to mitigate HOL blocking

24

HTTP/2: mitigating HOL blocking

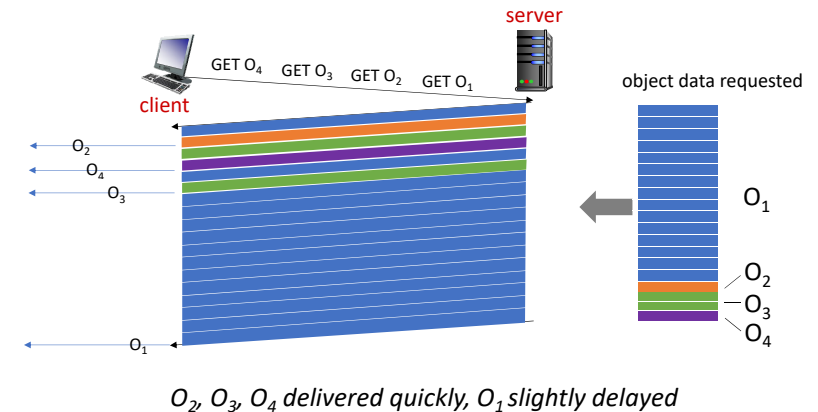
HTTP 1.1: client requests 1 large object (e.g., video file) and 3 smaller objects



25

HTTP/2: mitigating HOL blocking

HTTP/2: objects divided into frames, frame transmission interleaved



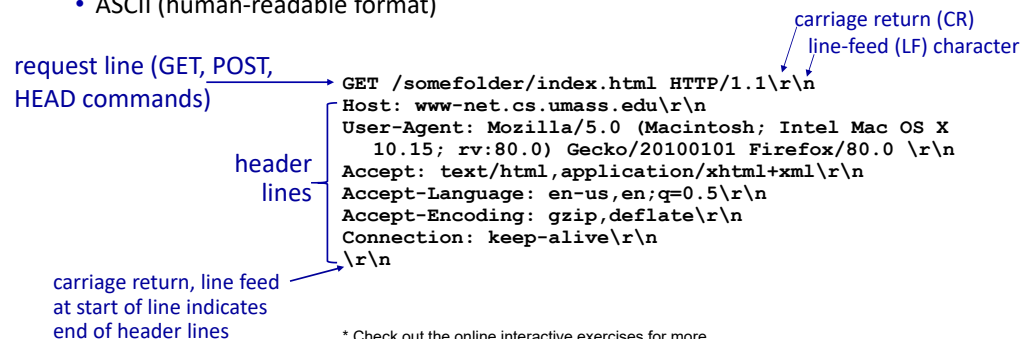
26

HTTP request message

two types of HTTP messages: *request*, *response*

HTTP request message:

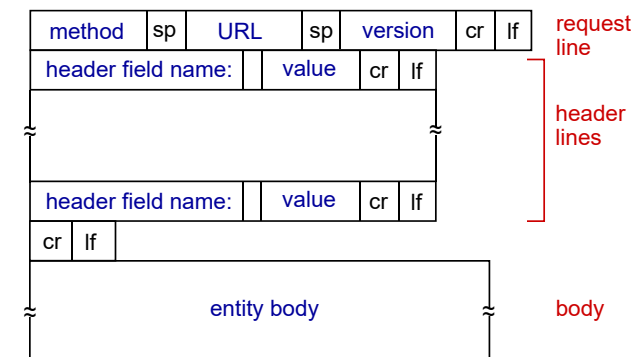
- ASCII (human-readable format)



* Check out the online interactive exercises for more examples: http://gaia.cs.umass.edu/kurose_ross/interactive/

27

HTTP request message: general format



28

Other HTTP request messages

POST method:

- web page might include form input or allow to upload files
- user input (sent from client to server) is put in entity body of HTTP POST request message

HEAD method:

- requests headers (only) that would be returned *if* specified URL were requested with an HTTP GET method.
- often for debugging

GET method (for sending data to server):

- include information in URL field of HTTP GET request message (following a '?'):

www.google.com/search?q=http+get&lr=lang_en

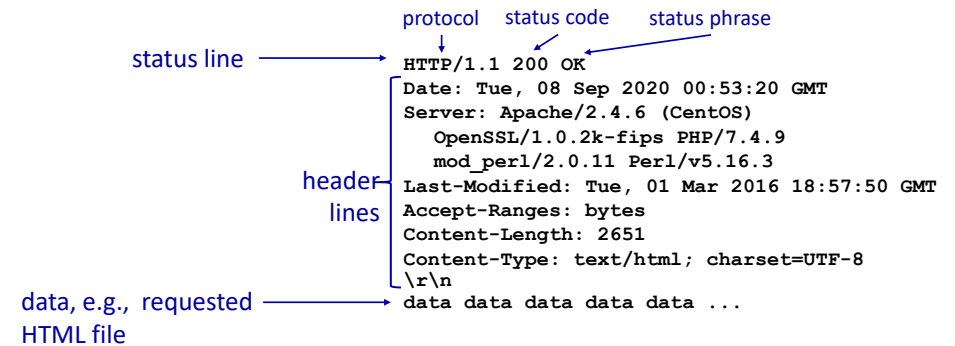
- entity body ignored or rejected

PUT method:

- create or update a file to server
 - with content in entity body
 - at specific URL

29

HTTP response message



* Check out the online interactive exercises for more examples:
http://gaia.cs.umass.edu/kurose_ross/interactive/

30

HTTP response status codes

- status code appears in 1st line in server-to-client response message.
- some sample codes:

200 OK

- request succeeded, requested object later in this message

301 Moved Permanently

- requested object moved, new location specified later in this message (in Location: field)

400 Bad Request

- request msg not understood by server

404 Not Found

- requested document not found on this server

505 HTTP Version Not Supported

31

Trying out HTTP (client side) for yourself

1. telnet to your favorite Web server:

```
% telnet www.cs.nthu.edu.tw 80
```

- opens TCP connection to port 80 (default HTTP server port) at www.cs.nthu.edu.tw.
- anything typed in will be sent to port 80 at www.cs.nthu.edu.tw

2. type in a GET HTTP request:

```
GET /~jungchuk/test.html HTTP/1.1
Host: www.cs.nthu.edu.tw
```

- by typing this in (**hit carriage return twice**), you send this minimal (but complete) GET request to HTTP server

3. look at response message sent by HTTP server!

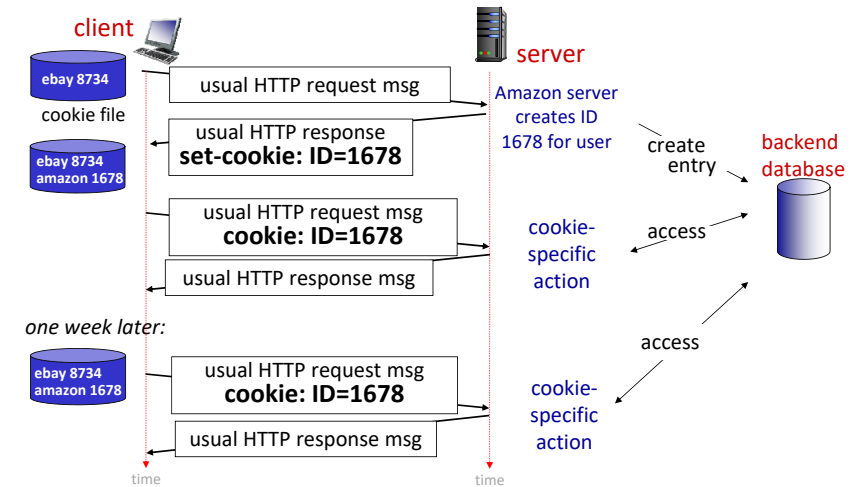
(or use Wireshark to look at captured HTTP request/response)

32

Maintaining user/server state: cookie


- **state** is very helpful for:
 - identification/authorization
 - shopping carts
 - recommendations
- Where to keep state?
 - at protocol endpoints
 - client side: browser
 - server side: backend database
 - in messages
 - HTTP messages carry state
- How exactly?
 - **cookie**
 - local storage
 - IndexedDB data
 - Session storage
 - ...
- Challenge: HTTP itself is **stateless**

Maintaining user/server state: cookie



Web cache (or called proxy)

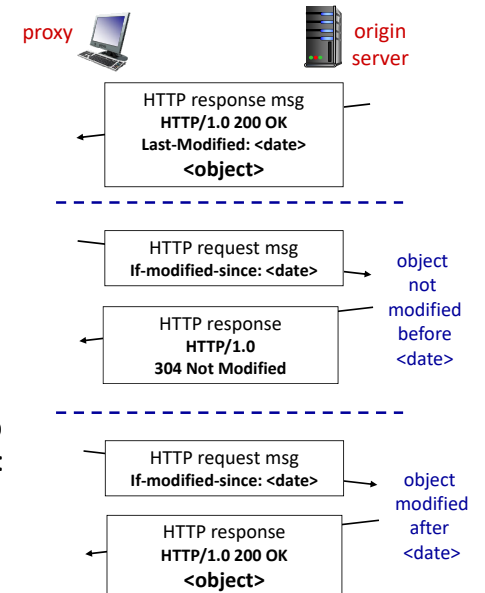
- **Goal:** satisfy client requests on behalf of origin server
 - user configures browser to point to a (local) **Web cache**
- Why web caching?
 - reduce response time for client request
 - cache is closer to client
 - reduce traffic on ISP's access link
 - enables "poor" content providers to more effectively deliver content
 - anonymity
- browser/client sends all HTTP requests to proxy
 - if object not in proxy: proxy requests object from origin server, caches received object, then returns object to client
 - else proxy returns object to client



Check freshness by conditional GET

Goal: don't send object if proxy has cached up-to-date version

- "If-modified-since" header
- "Expires" or "max-age" header
- **proxy:** specify date of cached copy in HTTP request
If-modified-since: <date>
- **origin server:** response contains no object if cached copy is up-to-date:
HTTP/1.0 304 Not Modified



Caching example


Scenario:

- access link rate: 15 Mbps
- rate of demand over access link/LAN = ?
 - web object size: 0.1 Mb
 - average request rate: 150 requests/sec
- one-way delay in public Internet: 2 seconds
- link/LAN delay is 10 ms at low or medium traffic
is minutes at high utilization

Performance:

- access link utilization = $15/15 = 1$
- LAN utilization: $15M/1G = 0.015$
- end-to-end delay
= Internet delay + access link delay + LAN delay
≈ 2 sec + minutes + 10 ms

access link is the bottleneck



37

Option 1: buy a faster access link


Scenario:

- access link rate: ~~15~~ 100 Mbps
- rate of demand over access link/LAN = 15 Mbps
 - web object size: 0.1 Mb
 - average request rate: 150 requests/sec
- one-way delay in public Internet: 2 seconds
- link/LAN delay is 10 ms at low or medium traffic
is minutes at high utilization

Performance:

- access link utilization = $15/100 = 0.15$
- LAN utilization: 0.015
- end-to-end delay
= Internet delay + access link delay + LAN delay
≈ 2 sec + minutes + 10 ms

Cost: faster access link (expensive!)



38

Option 2: install a proxy


Scenario:

- access link rate: 15 Mbps
- rate of demand over access link/LAN = 15 Mbps
 - web object size: 0.1 Mb
 - average request rate: 150 requests/sec
- one-way delay in public Internet: 2 seconds
- link/LAN delay is 10 ms at low or medium traffic
is minutes at high utilization

Cost: proxy (cheap!)

Performance:

- LAN utilization: .?
- access link utilization = ?
- average end-end delay = ?




39

Calculating end-to-end delay with proxy

suppose cache hit rate is 0.4:

- each request satisfied by proxy takes
≈ 10 ms = 0.01 s
- each request served by origin server takes
 - access link utilization = $0.6 * 15 / 15 = 0.6$
 - access link delay ≈ 10 ms
 - end-to-end delay
≈ 2 s + 10 ms + 10 ms ≈ 2.02 s
- average end-end delay:
= $0.6 * (\text{delay from origin server})$
+ $0.4 * (\text{delay when satisfied by proxy})$
≈ $0.6 * 2.02 \text{ s} + 0.4 * 0.01 \text{ s}$
≈ 1.2 secs

lower avg end-end delay than with 100 Mbps link (and cheaper too!)



40