

## Chapter 3: roadmap

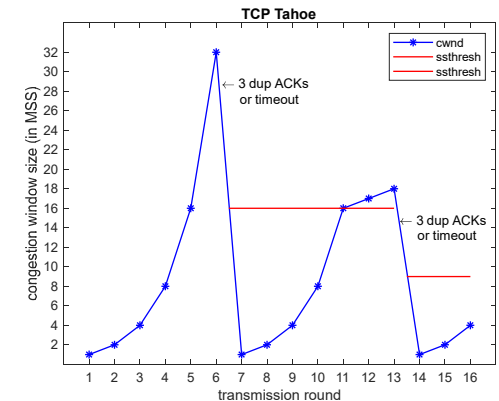
- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- Principles of reliable data transfer
- Connection-oriented transport: TCP
- Principles of congestion control
- TCP congestion control
- Evolution of transport-layer functionality



99


## TCP congestion control: overview

- sender adjusts *cwnd* dynamically to limit the tx rate ( $cwnd \propto \text{rate}$ )
- sender perceives congestion by
  - timeout and 3 duplicate ACKs
- TCP Tahoe:
  - “slow start” phase
    - start with *cwnd* = 1 (MSS)
    - if no loss occurs, double *cwnd* every RTT (until *cwnd* reaches *ssthresh* and goes to “congestion avoidance”)
    - if loss occurs, reset *cwnd* = 1
  - “congestion avoidance” phase
    - if no loss occurs, *cwnd* grows linearly
    - if loss occurs, go to “slow start”



100

## TCP congestion window: details




- TCP sender limits transmission:  $\text{LastByteSent} - \text{LastByteAcked} \leq cwnd$ 
  - *roughly*: send *cwnd* bytes, wait RTT for ACKs, then send more bytes

$$\text{TCP rate} \approx \frac{cwnd}{RTT} \text{ bytes/sec}$$

101

## TCP slow start: details

- when connection begins, increase rate exponentially until first loss event:
  - initially *cwnd* = 1 MSS
  - double *cwnd* every RTT
    - done by incrementing *cwnd* for every ACK received
- initial rate is slow, but ramps up exponentially fast



102

## Improvement for congestion avoidance: AIMD


- **approach:** senders can increase sending rate until packet loss (congestion) occurs, then decrease sending rate on loss event in a AIMD manner

### Additive Increase

increase sending rate by 1 maximum segment size every RTT until loss detected

### Multiplicative Decrease

cut sending rate in half at each loss event

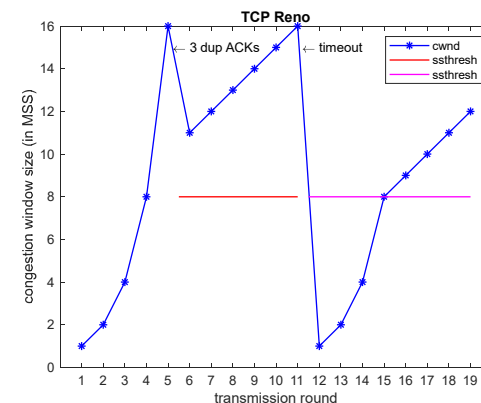


103

## TCP Reno

TCP Reno: cwnd is


- Cut to 1 MSS when loss detected by **timeout**
  - then enter "slow start"
  - $ssthresh \leftarrow cwnd/2$   
 $cwnd \leftarrow 1 \cdot MSS$
- Cut in half on loss detected by **3 duplicate ACK**
  - then enter "congestion avoidance"
  - more precisely,  
 $ssthresh \leftarrow cwnd/2$   
 $cwnd \leftarrow cwnd/2 + 3 \cdot MSS$



104

## TCP CUBIC (default in Linux)

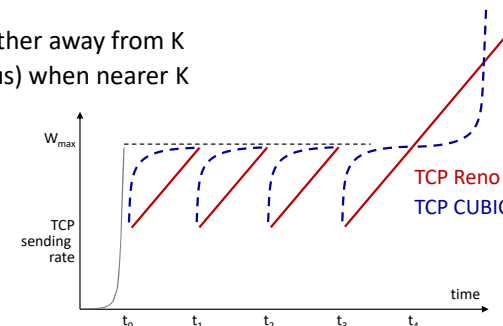
- Is there a better way than AIMD to "probe" for usable bandwidth?
- Insight/intuition:
  - $W_{max}$ : sending rate at which congestion loss was detected
  - congestion state of bottleneck link probably (?) hasn't changed much
  - after cutting rate/window in half on loss, initially ramp up to  $W_{max}$  *faster*, but then approach  $W_{max}$  more *slowly*



105

## TCP CUBIC


- K: point in time when TCP window size will reach  $W_{max}$ 
  - K itself is tuneable
- increase W as a function of the *cube* of the distance between current time and K
  - larger increases when further away from K
  - smaller increases (cautious) when nearer K
- TCP CUBIC default in Linux, most popular TCP for popular Web servers



106

## TCP fairness

**Fairness goal:** if  $K$  TCP sessions share same bottleneck link of bandwidth  $R$ , each should have average rate of  $R/K$




107

## Q: is TCP Reno Fair?

Example: two competing TCP connections:

- additive increase gives slope of 1, as throughput increases
- multiplicative decrease decreases throughput proportionally



**Is TCP fair?**

**A:** Yes, under idealized assumptions:

- same RTT
- fixed number of sessions only in congestion avoidance

108

## Fairness: must all network apps be “fair”?

### Fairness and UDP

- multimedia apps often do not use TCP
  - do not want rate throttled by congestion control
- instead use UDP:
  - send audio/video at constant rate, tolerate packet loss
- there is no “Internet police” policing use of congestion control

### Fairness, parallel TCP connections

- application can open *multiple* parallel connections between two hosts
- web browsers do this, e.g., link of rate  $R$  with 9 existing connections:
  - new app asks for 1 TCP gets rate  $1R/10$
  - new app asks for 11 TCPs, gets  $11R/20$

109